Instacart Agent-Readiness — Case Study Responses

Product Lead perspective, written in plain English

1. Why groceries are a good place to start

Groceries are boring in the best possible way. People buy roughly the same stuff every week — milk, bread, eggs, atta, dal. That means an AI assistant has plenty of past orders to learn from, and the user already knows what a 'good' order looks like when it arrives at their door.

The stakes per order are low. If the agent picks the wrong brand of bread, you eat different bread this week. Nobody loses a house, a flight, or their savings. Compare that to an agent booking flights or moving money — one mistake there is a disaster. Groceries let agents mess up, learn, and get better without hurting anyone badly.

The task is also naturally repeatable. Weekly cycles mean we get real feedback fast — 52 chances a year per user to improve, versus once-a-year purchases like insurance. And finally, substitutions ("no whole milk, get 2% instead") are the exact kind of judgment call agents are built for. Groceries are the training wheels for the whole agent economy.

2. The flashy demo vs. the durable feature

I would push back, but not with a flat 'no'. The board wants to see progress in six weeks; that's a real need, not a silly one. Killing the demo hurts the relationship and probably the budget.

So I'd offer a trade: we build a smaller, honest version of the flashy thing — enough to demo convincingly on stage — and spend the saved time on the foundation that will actually survive contact with real agents. In the room, I'd say plainly: 'If we ship the full flashy version, in three months you'll be on stage explaining why it broke. Let's show the board something real instead of something we'll regret.'

The job of a consultant isn't to say yes; it's to keep the client from a predictable mistake. I'd document my recommendation so if they overrule me, it's their call on record.

3. Checkout API vs. Substitution-preference API — pick one

I'd pick the substitution-preference API. Here's the plain-English reason: checkout is a solved problem. Every shopping app on earth already has a working checkout, and an agent can muddle through it. What agents can't do today is make the judgment call in the middle of a shop — 'they're out of Amul butter, is Britannia okay, or should I skip it?'

That's the moment where agents either feel like a trusted personal shopper or feel like a robot that keeps calling you. If we get substitutions right, the whole experience feels magical. If we only get checkout right, we've built a slightly better button.

Also, substitution preferences are where the user's trust boundaries live ('never swap my baby formula, always swap my snacks'). Getting that API right is what lets the agent act without waking the user up at 11pm to ask a question.

4. Two agents read 'quantity: 2' differently — where's the bug?

The bug is not in either agent. The bug is in our API. We shipped an ambiguous field and let each agent guess what it meant. That's our fault, not theirs.

The fix is to stop letting quantity mean two things. Split the field: 'units: 2' (I want 2 bottles) versus 'pack_size: 1kg, packs: 2' (I want 2 of the 1kg pack). Every product listing should say what a 'unit' is — a bottle, a pack, a kilo — so agents don't have to guess.

Long-term, this is why agent-readiness is a real discipline. Human users forgive ambiguity because they look at the picture and the price and figure it out. Agents can't. Every field has to mean exactly one thing. We'd also add a test suite that runs the top 5 agent platforms against our API and flags any place two of them disagree — that's our early-warning system.

5. Order threshold per user per month, and why

I'd start with a soft cap around 6 agent-placed orders per user per month — roughly a weekly shop plus one or two top-ups. Not because 6 is a magic number, but because it matches how real households actually buy groceries. Anything much above that and it's probably the agent misfiring, not the human genuinely needing more.

The cap should be soft, not hard. If a user hits it, the next order pauses for a quick confirmation ('your agent wants to place a 7th order this month, approve?') rather than getting blocked. That way we catch runaway agents without frustrating the family that genuinely orders twice a week.

The reason to be conservative early: we have zero historical data on agent-initiated purchases, and the first big 'my agent ordered ■15,000 of stuff I didn't want' story in the press will set the whole industry back a year. Start tight, watch the complaint rate, and loosen the cap once we have three to six months of clean data.

6. 'Should we be worried this happens to us?'

I'd answer honestly: yes, it could happen to us — and here's specifically why it probably won't.

Pretending we're immune is what got the other team in trouble. So I'd name the three things that most likely sank them — usually it's (a) building the flashy stuff before the foundation, (b) shipping ambiguous APIs that different agents interpret differently, and (c) no guardrails on agent behavior so one bad week of orders makes the news. Then I'd show them, concretely, what we're doing differently on each one.

I'd also tell them what I'm actually worried about on our project — the two or three real risks I'm tracking — and what would make me raise a flag. Clients don't lose trust when things go wrong; they lose trust when they find out things were going wrong and nobody told them. My job is to be the person who tells them early, not the person who gets caught out.

Bottom line: ‘Worry is fair. Panic isn’t. Here’s what we’re doing about it, and here’s what I’ll tell you the moment I see it going sideways.’